

Six questions to ask any AI-visibility number

The Invisible Funnel · Anton Andrusenko · SaaS Marketers Connect, Berlin, 30 Sep 2026 · every source: github.com/AntonAndrusenko/invisible-funnel

Ask them of a vendor's number, an agency's, your own team's, and mine.

1. **Where was it measured?** In the app or website people use, or through an API? For the same question, API and interface answers shared only about 12–15% of cited domains in one study ([Uberti-Bona Marin et al., 2026](#); not among the 51 cited in the talk).
2. **Which model and which plan?** ChatGPT's free and paid plans now run different models.
3. **Signed in or not?** Memory and personalisation change the answer.
4. **Which region and which language?** German and English are two measurements, not one.
5. **How many repeats?** A single run per question can't tell a real change from noise.
6. **Which denominator, and what was excluded?** All answers, answers naming any brand, or all mentions? Were failed runs and refusals dropped or counted as zero?

Definitions

- **Found · Cited · Recommended · Chosen:** four different outcomes (see the full pack).
- **Presence:** answers naming you ÷ all eligible answers, with an interval.
- **Share of mentions:** your mentions ÷ all brand mentions.
- **Prominence:** first / top three / later, reported as a distribution.
- **Accuracy:** each material claim coded as correct, outdated, omitted, wrong or non-answer.
- **Citation profile:** owned, competitor, media, directories, reviews, community.
- **Stability:** how much the brand set changes between repeats.
- **Business trace:** declared AI leads, AI-attributed leads, branded search and direct.

Same answers, three correct scores. 10 answers; 6 name at least one brand; you're in 5; 20 brand mentions in total → 50% of all answers, 83% of answers naming any brand, 25% of all mentions. Ask which one your dashboard shows.

A test you can run at B2B scale

Omio could test page groups across thousands of templated pages. A B2B site that gets about 1% of its visits from AI can't. So:

1. **Pick two clusters.** Change one topic cluster (for example, your pricing and integrations pages), and leave a similar cluster alone as the control.
2. **Measure both with the same fixed panel of buyer questions.** Use real questions from sales calls and support tickets, 8 runs each, before and after, for 4–6 weeks.
3. **Set guardrails.** Watch Search Console clicks on the changed pages (Google) and run an accuracy check on what the answers say (accuracy).
4. **Size it honestly.** With 30 questions you'll see only moves of about 20 points, so test big changes, not tweaks. Pilot 5 repeats per question first to learn your own noise.

How big a change you can detect (illustrative noise; your pilot decides):

Panel	Smallest change you can detect
30 questions × 8 runs	~20 points
50 × 8	~15 points
60 × 8	~14 points
~116 × 8	under 10 points
30 × 8, with my audit's noise on "listed first"	~31 points

My audit · 1 category, 1 brand · 10 questions × DE/EN · ChatGPT + Google AI Mode · 62 answers, signed-in sessions · 23 Sep 2026 · exploratory, not a trend

Extra runs per question buy little: extra questions buy more.

What to report instead of one score

- **Floor:** declared AI leads × close rate × average deal size.
- **Trend:** branded search and direct traffic, annotated with model releases and your own changes.
- **One holdout:** hold back the changes in one market or product line, and compare.

Five free fixes for Monday

1. **Check GA4's AI Assistant channel.** It has been built in since May 2026. Check which AI sources it catches, and add "an AI assistant" as an answer option to "How did you hear about us?".
2. **Check bot access, CDN settings and rendering.**
 - Let OAI-SearchBot, Claude-SearchBot, PerplexityBot, Googlebot and Bingbot in.
 - If you use Cloudflare, choose "Disallow AI Training", not "Block" (since 15 Sep 2026, "Block" also blocks Googlebot and Bingbot).
 - Make sure your key pages work without JavaScript.
 - Check that Search Console's new generative-AI setting is still on "include".
3. **Build a product-facts page:** pricing model, plans, integrations, limits, languages, compliance (GDPR, works council), who it's for (one size range, stated the same way everywhere), what it doesn't do, and the date it was last updated.
4. **Close the llms.txt ticket, for AI search.** If coding assistants read your developer docs, that's a different case.
5. **Ask AI your buyers' price and downside questions,** in German and in English, and write down what it gets wrong.

Illustrative sizing uses realistic but made-up noise; your own pilot decides. I work at Semrush (part of Adobe), which sells AI-visibility tools. Full evidence pack, with a rung and a source for every number: github.com/AntonAndrusenko/invisible-funnel