

The evaluation journey starts inside AI.



Sarah is a composite · her questions are real

**She never
clicked.**

**What will your
analytics record?**



She never clicked.

What will your analytics record?





Hallo!

Anton Andrusenko

Senior Product Manager at Semrush in AI Labs

 anton.andrusenko@semrush.com

 linkedin.com/in/anton-andrusenko/



The Invisible Funnel

What actually works in AI visibility —
and what the evidence says doesn't



Anton Andrusenko
Senior Product Manager
Semrush AI Labs





Every number tonight has a rung.

I work at Semrush, which sells AI-visibility tools. Its research sits on the same ladder.



152 sources in. 51 on stage.

Only 9 were field experiments on live platforms or sites.



THE GAP

What happened before
analytics could see the buyer?

The buyer sees AI.

Analytics sees the arrival.



94%

of B2B buyers report using
AI in the buying process

~1%

of website visits
arrive from AI assistants



We see the result — not the decision.

–39.8%

outbound organic clicks when an AI Overview appeared

 R1 · field RCT



Agarwal & Sen · randomised field experiment · N=1,065 · US · Google · Jan–Feb 2026 · conditional: only when an AI Overview appeared

Five meanings hide inside “visibility”.



Your pages get found and cited. Your brand gets mentioned, recommended, chosen.



HOW AI PICKS

What happens after Sarah asks one simple question?





One question. Many queries.

Answer the detailed questions, not just head terms.



Which HR system fits
a company like ours?



HR software Germany, 150 staff

DATEV payroll integration

works council features



10.7

sub-queries per prompt ·
95% with zero measured
search volume



Answer the detailed questions, not just head terms.

R4 · vendor study



**Your site is
one voice
in the answer.**

11.6%

of ChatGPT's citations went to
the recommended tool's own site





German-language pages




English + other languages

R4 · VENDOR STUDY

German prompts don't guarantee German sources.

Where the sources on German
prompts come from



ENGLISH ALONE · AVERAGE ACROSS
ALL 6 LANGUAGES TESTED

9%

15%


24%

40%

Many results. A few passages. One answer.

Not one model — a pipeline each vendor runs differently.

VENDOR DOCS + MEASUREMENTS



What wins the pick: relevance to the exact question — and where the text sits.

LAB · review of 45 studies

WHAT HOLDS UP?

Which AI-visibility claims
survive contact with evidence?



Which move has the strongest experimental support?



AUDIENCE VOTE





The famous **+40%** was
measured **after the page**
was found.

+40%

share of words in the answer ·
5 pages already in the model's hands

The Princeton GEO Paper in Plain
English: 5 Tactics That Boost AI
Citation by **40%**

derivatex.agency · Apr 2026

9 GEO Tactics That Boost AI
Citation Rates by **40%**

xseek.io · Jul 2026

Generative Engine Optimization (GEO):
The Definitive Guide [2026]

geoptie.com

Princeton research shows GEO can boost AI visibility by **40%**

Headlines as published · checked 29 Sep 2026





The famous **+40%** was
measured **after the page**
was found.

+40%

share of words in the answer ·
5 pages already in the model's hands



Strongest support?

D — narrowly, and only in the lab.



No tactic is proven across sites and platforms.

C-SEO Bench (NeurIPS 2025) • SAGEO Arena (KDD 2026) • Ahrefs: schema test, 1,885 pages; lms.txt logs, 137,000 domains, May 2026; 75,000-brand correlations

WHAT ACTUALLY WORKS

If the hacks don't survive
end to end, what can we control?



**Boring
beats
magical.**





Be retrievable.



01

Access

Let search bots in

02

Rendering

Key content in
HTML, not JavaScript

03

Answer first

Title and first lines
answer the question



When search lists your site, your brand gets named far more often.

A big difference — but correlation, not cause.




When the brand's own site was among the AI's sources, the brand was named in **about half the answers**.

When neither the brand nor its site came up in search: **under 4%**.

Named = mentioned, not necessarily recommended. Correlation, not cause.

Same prompt, repeated runs:
+40 to +49 points



In most AI answers, your page is a title and a passage.

Write the title and first two lines as the answer.



CHATGPT CAPTURE

≈200 characters
per result

usually starting at the H1



OPENED PAGES

1.2% of retrieved pages
were opened

when opened, pages were read
in whole



GEMINI

≈2,000 words
of grounding per question


shared by all sources

Be explicit — or AI borrows someone else's numbers.



Written on the page

DATEV integration →
no answer got it wrong



Missing from the page

Pricing → 6 answers used
third-party prices, 2 credited
them to the vendor



Be present **where the machine reads.**



ChatGPT · 36 answers

22%

Own site

63%

competitor sites



Google AI Mode · 26 answers

8%

Own site

38%

blogs, comparison sites & directories



My audit · 1 category, 1 brand · 10 questions × DE/EN · ChatGPT + Google AI Mode · 62 answers, signed-in sessions · 23 Sep 2026 · exploratory, not a trend · share of cited domains

Mentions travel with AI visibility. So does brand size.

BET E · R4 · correlation

Brand mentions · **0.66**

Backlinks · **0.22**

Correlation, not proof. No controlled test
of an off-site lever has been published.

Ahrefs · 75,000 brands · Spearman correlation with AI visibility · 2025–2026



Big brands have more mentions, more searches and
more AI visibility, all at once. Nobody has separated them.



R1 · field test

Test it like CRO, not like SEO.

One change · one control ·
a Google guardrail.



 NOT SHIPPED

WHAT CAN WE MEASURE?

How much of the invisible funnel
can we actually observe?



What you can see today: reach, not revenue.



GA4 · AI Assistant
channel · built in
since May 2026



Search Console ·
AI impressions only



Bing Webmaster Tools ·
citations and grounding
queries



Server logs ·
model page reads



“How did you hear
about us?”



Checked 30 Sep 2026





R4 • correlation



Who gets the credit?



Report a floor, a trend and one holdout.



Conductor, 13,770 domains • one B2B firm, 213 leads (it sells AEO) • Similarweb B2C panel, ChatGPT, US desktop, Jul-Dec 2025

Same observations. **Three correct scores.**

20 brand mentions in total.



Same observations. **Three correct scores.**

20 brand mentions in total.



The dashboard answers a question someone chose for you.





Documentation review

Most tools don't publish how they count.

7

publish a formula


2

say what happens
to failed runs

0

disclose their
sentiment model

of 20 AI-visibility vendors





One number hides what moved.

98%

present · 51 of 52 unbranded answers (95% interval 90–100%)

≤16%

share of mentions (upper bound)



In 51 of 52 photos — but one face in six.

REPORT A VECTOR, NOT A SCORE · VENDOR A

PRESENCE

98%

= mentioned

PROMINENCE

1st on price:
1 of 7

≈ recommended

ACCURACY

3 wrong or
outdated facts

SEGMENT FIT

not measured

CITATION PROFILE

own site: 22%
ChatGPT · 8% AI Mode
= cited

STABILITY

same brands
in 3 of 22 repeats

BUSINESS TRACE

empty, on purpose
= chosen



MONDAY

What should I actually do
when I get back to work?



Monday morning: five free fixes.

Turn visibility into a repeatable process.

THREE MOVES

01



INFLUENCE

Retrievable · Explicit & correct ·
Present where machines read

02



MEASURE

Found · Cited · Recommended · Chosen

03



TEST

One change · one control · one guardrail

FIVE FREE FIXES

1



Check GA4's AI Assistant channel

2



Check bot access, CDN settings and rendering

3



Build a product-facts page

4



Close the llms.txt ticket

5



Ask AI your buyers' price and downside
questions, in German and English



**Don't optimise for a
visibility score.**

**Optimise to be found,
trusted and correctly
represented.**



Q&A



Documents



LinkedIn